P05

Inference of multicast routing trees and bottleneck bandwidths using end-to-end measurements

Authors	Sylvia Ratnasamy, Steven McCanne

Published	IEEE Conference on Computer Communications (INFOCOM) (1999)
DOI	https://doi.org/10.1109/INFCOM.1999.749302
Read it	https://doi.org/10.1109/INFCOM.1999.749302 - Subscription access: abstract and DOI provided; use your library login or ask the organisers
Domain	Networking & distributed systems

Abstract (quoted verbatim from the publication)

The efficacy of end-to-end multicast transport protocols depends critically upon their ability to scale efficiently to a large number of receivers. Several research multicast protocols attempt to achieve this high scalability by identifying sets of co-located receivers in order to enhance loss recovery, congestion control and so forth. A number of these schemes could be enhanced and simplified by some level of explicit knowledge of the topology of the multicast distribution tree, the value of the bottleneck bandwidth along the path between the source and each individual receiver and the approximate location of the bottlenecks in the tree. In this paper, we explore the problem of inferring the internal structure of a multicast distribution tree using only observations made at the end hosts. By noting correlations of loss patterns across the receiver set and by measuring how the network perturbs the fine-grained timing structure of the packets sent from the source, we can determine both the underlying multicast tree structure as well as the bottleneck bandwidths. The simulations show that the algorithm is robust and appears to converge to the correct tree with high probability.

Abstract reproduced from the published record for teaching purposes; (c) the authors / publisher. Source: https://doi.org/10.1109/INFCOM.1999.749302

Notes: record the page or section for every claim you make about this paper.

Reading prompts

1. What problem does the paper try to solve? Write it in one sentence, in your own words.
2. What is unknown or hidden, and what can actually be observed or measured?
3. What method do the authors use? Is it new, borrowed from another paper, or a modification of one?
4. What evidence do they give (proofs, simulations, experiments, real data)?
5. What assumptions or limitations can you find?
6. Which other papers in the corpus does it cite, use, extend or resemble? Note the page or section.

This card is a workshop wrapper prepared by the organisers. Only the abstract is quoted from the publication; read the paper itself through the link above before asserting any relationship.