P06

Learning Latent Tree Graphical Models

Authors Myung Jin Choi, Vincent Y. F. Tan, Animashree Anandkumar, Alan S. Willsky

Published Journal of Machine Learning Research (2011)

DOI https://doi.org/10.48550/arXiv.1009.2722

Read it https://arxiv.org/abs/1009.2722 - Free to read (follow the open-access link)

Domain Computer science, graph theory & machine learning

Abstract (quoted verbatim from the publication)

We study the problem of learning a latent tree graphical model where samples are available only from a subset of variables. We propose two consistent and computationally efficient algorithms for learning minimal latent trees, that is, trees without any redundant hidden nodes. Unlike many existing methods, the observed nodes (or variables) are not constrained to be leaf nodes. Our first algorithm, recursive grouping, builds the latent tree recursively by identifying sibling groups using so-called information distances. One of the main contributions of this work is our second algorithm, which we refer to as CLGrouping. CLGrouping starts with a pre-processing procedure in which a tree over the observed variables is constructed. This global step groups the observed nodes that are likely to be close to each other in the true latent tree, thereby guiding subsequent recursive grouping (or equivalent procedures) on much smaller subsets of variables. This results in more accurate and efficient learning of latent trees. We also present regularized versions of our algorithms that learn latent tree approximations of arbitrary distributions. We compare the proposed algorithms to other methods by performing extensive numerical experiments on various latent tree graphical models such as hidden Markov models and star graphs. In addition, we demonstrate the applicability of our methods on real-world datasets by modeling the dependency structure of monthly stock returns in the S&P index and of the words in the 20 newsgroups dataset.

Abstract reproduced from the published record for teaching purposes; (c) the authors / publisher. Source: https://doi.org/10.48550/arXiv.1009.2722

Notes: record the page or section for every claim you make about this paper.

Reading prompts

1. What problem does the paper try to solve? Write it in one sentence, in your own words.
2. What is unknown or hidden, and what can actually be observed or measured?
3. What method do the authors use? Is it new, borrowed from another paper, or a modification of one?
4. What evidence do they give (proofs, simulations, experiments, real data)?
5. What assumptions or limitations can you find?
6. Which other papers in the corpus does it cite, use, extend or resemble? Note the page or section.

This card is a workshop wrapper prepared by the organisers. Only the abstract is quoted from the publication; read the paper itself through the link above before asserting any relationship.