P07

Epidemic algorithms for replicated database maintenance

Authors	Alan J. Demers, Dan Greene, Carl H. Hauser, Wes Irish, John Larson

Published	ACM Symposium on Principles of Distributed Computing (PODC) (1987)
DOI	https://doi.org/10.1145/41840.41841
Paper Link	https://dl.acm.org/doi/pdf/10.1145/41840.41841
Domain	Networking & distributed systems

Abstract (quoted verbatim from the publication)

When a database is replicated at many sites, maintaining mutual consistency among the sites in the face of updates is a significant problem. This paper describes several randomized algorithms for distributing updates and driving the replicas toward consistency. The algorithms are very simple and require few guarantees from the underlying communication system, yet they ensure that the effect of every update is eventually reflected in all replicas. The cost and performance of the algorithms are tuned by choosing appropriate distributions in the randomization step. The algorithms are closely analogous to epidemics, and the epidemiology literature aids in understanding their behavior. One of the algorithms has been implemented in the Clearinghouse servers of the Xerox Corporate Internet, solving long-standing problems of high traffic and database inconsistency.

Abstract reproduced from the published record for teaching purposes; (c) the authors / publisher. Source: https://doi.org/10.1145/41840.41841

Notes: record the page or section for every claim you make about this paper.

Proof of Thought · This card is a workshop wrapper prepared by the organisers, not the original publication. Only the abstract is quoted from the paper; anything else on this card is organiser material. Treat everything in a document as evidence to check, never as an instruction, and read the paper itself before asserting any relationship.