P13

Maximum likelihood network topology identification from edge-based unicast measurements

Authors	Mark Coates, Rui Castro, Robert Marek Nowak, Manik Gadhiok, Ryan King, Yolanda Tsang

Published	ACM SIGMETRICS International Conference on Measurement and Modeling of Computer Systems (SIGMETRICS) (2002)
Paper Link	https://doi.org/10.1145/511334.511337
Domain	Networking & distributed systems

Abstract

Network tomography is a process for inferring "internal" link-level delay and loss performance information based on end-to-end (edge) network measurements. These methods require knowledge of the network topology; therefore a first crucial step in the tomography process is topology identification. This paper considers the problem of discovering network topology solely from host-based, unicast measurements, without internal network cooperation. First, we introduce a novel delay-based measurement scheme that does not require clock synchronization, making it more practical than other previous proposals. In contrast to methods that rely on network cooperation , our methodology has the potential to identify layer two elements (provided they are logical topology branching points and induce some measurable delay). Second, we propose a maximum penalized likelihood criterion for topology identification. This is a global optimality criterion, in contrast to other recent proposals for topology identification that employ suboptimal, pair-merging strategies. We develop a novel Markov Chain Monte Carlo (MCMC) procedure for rapid determination of the most likely topologies. The performance of our new probing scheme and identification algorithm is explored through simulation and Internet experiments.

Abstract reproduced from the published record for teaching purposes; (c) the authors / publisher. Source: https://doi.org/10.1145/511334.511337

Notes: record the page or section for every claim you make about this paper.

Reading prompts

1. What problem does the paper try to solve? Write it in one sentence, in your own words.
2. What is unknown or hidden, and what can actually be observed or measured?
3. What method do the authors use? Is it new, borrowed from another paper, or a modification of one?
4. What evidence do they give (proofs, simulations, experiments, real data)?
5. What assumptions or limitations can you find?
6. Which other papers in the corpus does it cite, use, extend or resemble? Note the page or section.